(/)

LOGOUT

OPEN HACK ENVIRONMENT OVERVIEW OPEN HACK GUIDE

PROVIDE FEEDBACK

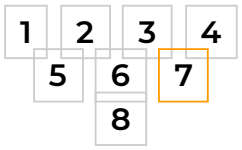

Mark Complete

Challenge 7: Safety in Depth

Background

The Adventure Works data science team wants to experiment with creating a custom object detection model using a convolutional neural network (CNN). This custom model should perform the same object detection task as the model previously built with the *Computer Vision* service, detecting whether or not each person in an image is protected by a helmet.

Prerequisites

- An environment for sharing code and working in Jupyter.
- An installation of a deep learning framework with which to train an object detection model.

(/)

LOGOUT

- The safety dataset you downloaded and extracted in the previous challenge from https://computervisionhack.blob.core.windows.net/challengefiles/safety_images.zip (https://computervisionhack.blob.core.windows.net/challengefiles/safety_images.zip

Challenge

1. Using the deep learning framework of your choice, create an object detection solution. This model should be able to detect and create a bounding box around each helmet present in an image.
2. Test your model using an image that was not included in the training dataset, displaying the detected objects, their classes, and their bounding boxes.

Hints

- The VoTT utility can output tags for training in formats suitable for the commonly used deep learning frameworks.
- There are many existing models for object detection, such a *YOLO*, *FAST R-CNN* and *Faster R-CNN* that you can use as the base model for *transfer learning*.

Success Criteria

- Your object detection model must achieve a Mean Average Precision (MAP) of **60%** or higher.
- You must write code that uses the model to get predictions for the classes and locations of objects in a test image, and plots the image overlaid with annotated bounding boxes for the predicted classes, like this:

(/)

Helmet Detection

References

Concepts

- What is object detection? (https://tryolabs.com/blog/2017/08/30/object-detection-an-overview-in-the-age-of-deep-learning/)
- What is MAP? (http://fastml.com/what-you-wanted-to-know-about-mean-average-precision/)

Tools and Frameworks

- Visual Object Tagging Tool (VOTT) (https://github.com/Microsoft/VoTT)
- Faster R-CNN in PyTorch (https://github.com/ruotianluo/pytorch-faster-rcnn)
- The Tensorflow Object Detection API (https://github.com/tensorflow/models/tree/master/research/object_detection)
- Object Detection with CNTK (https://docs.microsoft.com/en-us/cognitive-toolkit/Object-Detection-using-Faster-R-CNN)

(/)

LOGOUT